

Special Communication

Deep learning for electronic health records: A comparative review of multiple deep neural architectures



Jose Roberto Ayala Solares^{a,b}, Francesca Elisa Diletta Raimondi^a, Yajie Zhu^{a,*},
Fatemeh Rahimian^a, Dexter Canoy^{a,b,c}, Jenny Tran^a, Ana Catarina Pinho Gomes^a,
Amir H. Payberah^a, Mariagrazia Zottoli^a, Milad Nazarzadeh^{a,d}, Nathalie Conrad^a,
Kazem Rahimi^{a,b}, Gholamreza Salimi-Khorshidi^a

^a The George Institute for Global Health (UK), University of Oxford, United Kingdom

^b NIHR Oxford Biomedical Research Centre, Oxford University Hospitals NHS Foundation Trust, United Kingdom

^c Faculty of Medicine, University of New South Wales, Sydney, Australia

^d Collaboration Center of Meta-Analysis Research, Torbat Heydariyeh University of Medical Sciences, Torbat Heydariyeh, Iran

ARTICLE INFO

Keywords:

Deep learning
Representation learning
Neural networks
Electronic health records
CPRD

ABSTRACT

Despite the recent developments in deep learning models, their applications in clinical decision-support systems have been very limited. Recent digitalisation of health records, however, has provided a great platform for the assessment of the usability of such techniques in healthcare. As a result, the field is starting to see a growing number of research papers that employ deep learning on electronic health records (EHR) for personalised prediction of risks and health trajectories. While this can be a promising trend, vast paper-to-paper variability (from data sources and models they use to the clinical questions they attempt to answer) have hampered the field's ability to simply compare and contrast such models for a given application of interest. Thus, in this paper, we aim to provide a comparative review of the key deep learning architectures that have been applied to EHR data. Furthermore, we also aim to: (1) introduce and use one of the world's largest and most complex linked primary care EHR datasets (i.e., Clinical Practice Research Datalink, or CPRD) as a new asset for training such data-hungry models; (2) provide a guideline for working with EHR data for deep learning; (3) share some of the best practices for assessing the "goodness" of deep-learning models in clinical risk prediction; (4) and propose future research ideas for making deep learning models more suitable for the EHR data. Our results highlight the difficulties of working with highly imbalanced datasets, and show that sequential deep learning architectures such as RNN may be more suitable to deal with the temporal nature of EHR.

1. Introduction

Electronic health records (EHR) systems store data associated with each individual's health journey (including demographic information, diagnoses, medications, laboratory tests and results, medical images, clinical notes, and more) [1]. While the primary use of EHR was to improve the efficiency and ease of access of health systems [2], it has found a lot of applications in clinical informatics and epidemiology [3–5]. In particular, EHR have been used for medical concept extraction [6,7], disease and patient clustering [8,9], patient trajectory modelling [10], disease prediction [11,12], and data-driven clinical decision support [13,14], to name a few.

The early analyses of EHR relied on simpler and more traditional statistical techniques [15]. More recently, however, statistical machine

learning techniques such as logistic regression [16], support vector machines (SVM) [17], Cox proportional hazard model [18], and random forest [19] have also been employed for mining reliable predictive patterns in EHR data. While the simplicity and interpretability of such statistical models are desirable for medical applications, their weakness in dealing with high-dimensional input, their reliance on many assumptions, both statistical and structural, and their need for hand-crafted features/markers (guided by a domain expertise) make their use for comprehensive analyses of EHR data impractical [20–22]. To alleviate these issues, one needs to analyse each individual's entire medical history (i.e., a series mixed-type and multimodal data packed in irregular intervals) [22], using modelling techniques that can discover and take into account complex nonlinear interactions among variables [23,24].

* Corresponding author.

E-mail address: yajie.zhu@georgeinstitute.ox.ac.uk (Y. Zhu).

<https://doi.org/10.1016/j.jbi.2019.103337>

Received 4 March 2019; Received in revised form 25 September 2019; Accepted 4 November 2019

1532-0464/ © 2019 Elsevier Inc. All rights reserved.




Fig. 1. Published studies found on Semantic Scholar (<https://www.semanticscholar.org/>) through October 2018 using keywords “deep learning” AND (“electronic health records” OR “electronic medical records”) in the title or abstract.

The past decade has witnessed the profound impact of DL in a broad range of applications, including but not limited to its striking performance in natural language processing [25], speech recognition [26], computer vision [27], and reinforcement learning [28]. As a result, in recent years, the field of health informatics has seen a growth in the use of DL for a broad range of tasks [29]. For instance, DL has been extensively applied in the analysis of medical images [30], with Google’s work on detection of diabetic retinopathy [31], and detection of cancer metastases on gigapixel pathology images [32], being the most widely known examples. In genomics, DL has helped identify common genetic variations more accurately than conventional methods [33,34], opening the doors to new directions in genomics research [35,36] and drug discovery for personalised medicine [37]. In all these applications, DL algorithms analysed big datasets to produce a compact general-purpose set of features that can be used for personalised risk prediction, treatment recommendations, and clinical trial recruitment [38,39].

The number of research projects using DL for the analysis of EHR has been growing rapidly in the past few years (see Fig. 1). A quick search will show that the key DL architectures (e.g., feed-forward neural networks (FFNN), convolutional neural networks (CNN), and recurrent neural networks (RNN)), have been employed for the analysis and modelling of EHR. In one of the earliest such works, Tran et al. [40] introduced eNRBM (electronic medical records-driven nonnegative restricted Boltzmann machines), for learning a universal representation from complete EHR data (i.e., an automatic feature extraction). eNRBM introduced constraints for nonnegative coefficients (for explainability of the learned parameters) and structural smoothness (to account for disease hierarchy in classification coding schemes) to standard Bernoulli RBM; the resulting model showed a superior performance for suicide risk prediction, when compared with “manual feature engineering”. Furthermore, eNRBM was shown to be capable of providing clinically meaningful clustering of patients and diseases. Similarly, Miotto et al. [38] chose deep stacked denoising autoencoder (SDA) for training a universal feature extractor, and showed that it outperforms the expert-driven feature engineering in a variety of clinical risk-prediction tasks, including diabetes mellitus with complications, cancer of rectum and anus, cancer of liver and intrahepatic bile duct, congestive heart failure (non-hypertensive), among others. In both approaches, a supervised learning model (such as logistic regression, SVM, or random forests) was needed to map the learned representation to the outcome of interest. For this reason, we refer to such approaches as “modular”.

Neither eNRBM nor SDA considered the temporal information in the EHR, explicitly; the attempt by Tran et al. [40] to split patients’ journeys into non-overlapping time intervals, and concatenating them (see the details in Section 2) was the closest that modular models have become to considering time. In order to address this shortcoming, Nguyen

et al. [20] introduced a CNN architecture called Deepr (Deep record), where a patient’s journey is modelled as a long sentence of medical codes, with each code embedded into a new space to facilitate the statistical and algebraic operations (i.e., similar to word embedding in natural language processing [41]), and denoting the time between events as “special words”. Deepr is an example of an “end-to-end” model, trained to map the EHR history directly to the outcome of interest. Such a model was validated on hospital data to predict unplanned readmission after discharge, obtaining a slightly better performance with an AUC of 0.80 (3-month prediction) and 0.819 (6-month prediction), compared to logistic regression with bag-of-words representation, where an AUC of 0.797 (3-month prediction) and 0.811 (6-month prediction) was obtained.

In another end-to-end modelling endeavour, Choi et al. introduced Med2Vec [42], an FFNN model for learning representations for both visits (a group of medical codes) and medical codes, providing word embeddings comparable to other techniques like Skip-gram [43], GloVe [44], and stacked autoencoders [38]. Furthermore, in multiple studies [45–47], Choi et al. extended their work by using an RNN architecture to detect influential past visits and significant clinical variables while remaining clinically interpretable. In particular, the RETAIN (REverse Time Attention) model [45] used two RNNs and a two-level neural attention model to process sequential data using an end-to-end approach. The model obtained an AUC of 0.8705 when tested on heart failure diagnosis from Sutter Health medical records, suggesting that the model had a high predictive accuracy, while providing interpretability in the results. Since its development, the RETAIN model has been enhanced considerably. For instance, Kwon et al. [48] produced an interactive visual interface named RetainVis, that offers insights into how individual medical codes contribute to making risk predictions, Ma et al. [49] developed an attention-based bidirectional RNN that measures the relationships of past and future visits (from Medicaid claims in the US) to produce diagnosis prediction, while providing clinically meaningful interpretations, and Choi et al. [50] built a graph-based attention model, capable of using hierarchical information inherent to medical ontologies, resulting in a 10% higher accuracy for predicting rare diseases and 3% improved AUC for predicting heart failure using an order of magnitude less training data.

RNN’s power in modelling language and other sequences [51,52], when paired with their early success in learning new representations of EHR (as introduced above) made it a popular choice for DL researchers in this domain. Pham et al. [21] introduced an RNN architecture (called DeepCare) to predict future medical outcomes based on the patients’ health trajectories. To do so, they used a modified long short-term memory (LSTM) unit [53,54] with a forgetting mechanism and the ability to handle irregular inter-visit intervals. The model was applied to prediction of unplanned readmission within 12 months, where an F-score of 0.791 was obtained, an improvement over traditional machine learning techniques like SVM (F-score of 0.667) and random forests (F-score of 0.714). Extending this idea, Rajkomar et al. [22] introduced an ensemble model [55], which combined the strengths of three different models: a weighted RNN, a feedforward model with time-aware attention, and a boosted embedded time series model. The authors applied their model to EHR data from the University of California, San Francisco, and the University of Chicago Medicine. They demonstrated the effectiveness of this DL model in a variety of clinical tasks, including mortality prediction (AUC of 0.93–0.95), unexpected readmission (AUC of 0.76–0.77), and increased length of stay (AUC of 0.85–0.86). In addition, Ma et al. [56] developed a hybrid network structure composed of both an RNN and a CNN in order to extract comprehensive multifaceted patient information patterns with attentive and time-aware modulators. Such a network was applied in the prediction of heart failure (SNOW dataset) and diabetes mellitus (EMRBots dataset), obtaining AUCs of 0.729 (heart failure) and 0.943 (diabetes mellitus), outperforming other models like logistic regression (0.650/0.790) and RETAIN (0.668/0.767).

Table 1
Main approaches of Deep Learning for EHR (the abbreviations' meaning can be found in [Appendix A](#)).

Model	Architecture	Learning Process	# Patients considered	Predictors Considered	Outcome	Performance Metrics
eNRBM [40]	AE	Modular	7,578	Diagnoses (ICD-10), procedures (ACHD), Elixhauser comorbidities, diagnosis related groups, emergency attendances and admissions, demographic variables (ages in 10 year intervals and gender)	suicide risk prediction	F-score Recall Precision
Deep Patient [38]	SDA	Modular	704,587	Demographic variables (i.e., age, gender, and race), diagnoses (ICD-9), medications, procedures, lab tests, free-text clinical notes	future disease prediction	AUROC Accuracy F-score Precision AUROC
DeepR [20]	CNN	End-to-end	300,000	Diagnoses (ACS), procedures (ACHD)	unplanned readmission prediction	Precision F-score Recall
DeepCare [21]	RNN	End-to-end	7,191	Diagnoses (ICD-10), procedures (ACHD), medications (ATC)	disease progression, unplanned readmission prediction	
Med2Vec [42]	FFNN	End-to-end	ND	Diagnoses (ICD-9), procedures (CPT), medications (NDC)	medical codes in previous/future visits	AUROC Recall
Doctor AI [46]	RNN	End-to-end	263,706	Diagnoses (ICD-9), procedures (CPT), medications (GPI)	medical codes in future visits, duration until next visit	Recall R^2
RETAIN [45]	RNN	End-to-end	32,787	Diagnoses (ICD-9), procedures (CPT), medications (GPI)	heart failure prediction	AUROC
RetainVis [48]	RNN	End-to-end	63,030 (heart failure) 117,612 (cataract)	Diagnoses (ICD-9), procedures, medications	heart failure prediction, cataract prediction	AUROC AUPRC
Ensemble model [22]	RNN/FFNN	End-to-end	216,221	Demographics, provider orders, diagnoses, procedures, medications, laboratory values, vital signs, flowsheet data, free-text medical notes	inpatient mortality, 30-day unplanned readmission, long length of stay, diagnoses	AUROC

Despite the growth in the number of DL papers for EHR, the majority of these papers have used data from different countries and healthcare systems, employed different data cleaning and preparation processes (e.g., inclusion/exclusion criteria, input fields, and sample size), developed and used different DL architectures, while answering different clinical questions; even when assessing the models for risk prediction, the reviewed papers were not consistent in the quantification of model performance (see [Table 1](#)). Such wide paper-to-paper variability will hamper the ability of the field to compare and contrast the relevant papers for a clinical question at hand.

This paper primarily aims to provide a comparative review of the key DL architectures that have been used for the analysis of EHR data; describing their corresponding strengths and weaknesses for various real-world challenges when applied to a single complex EHR. We focused our work on getting an efficient patient representation to tackle two important but diverse clinical prediction challenges: predicting emergency admission, or heart failure. In addition to being the first comparative review of key DL architectures for EHR data, this paper also aims to: (1) introduce and use one of the world's largest and most complex databases of linked primary care EHR (i.e., Clinical Practice Research Datalink, or CPRD) that captures 30 years of medical history of patients, as a new asset for training such data-hungry models; (2) review some of the approaches for working with such data for DL; (3) share some of the best practices for assessing the performance of DL models in clinical risk prediction tasks; (4) and propose ideas of future research for making DL models even more suitable for EHR data.

2. Materials and methods

2.1. Overview

Our comparative review assessed four main DL architectures that have been employed so far for the analysis of EHR data: AE (auto-encoders), SDA, CNN, and RNN (for complete reference, please see [Table 1](#) with the corresponding abbreviations available in [Appendix A](#)). In particular, we used the following models: eNRBM [40], Deep Patient [38], DeepR [20], and RETAIN [45]; this is the best representative subset for all the key DL architectures in the field [3–5].

2.2. Clinical Practice Research Datalink (CPRD)

Over 98% of the UK population are registered with a general practice (GP), the first point of contact for healthcare in the UK National Health Service [57]. The CPRD is a service that collects de-identified longitudinal primary care data from a network of GPs in the UK, which are linked to secondary care and other health and area-based administrative databases [58]. Among these linked databases include the Hospital Episode Statistics (for data on hospitalisations, outpatient visits, accident and emergency attendances, and diagnostic imaging), the Office of National Statistics (death registration), Public Health England (cancer registration), and the Index of Multiple Deprivation. Around 1 in 10 GP units contribute data to the CPRD. To date, the CPRD covers 35 million patient lives among whom 10 million were currently registered patients from 674 GP units, making it one of the largest primary care EHR databases in the world.

Patients included in the CPRD database were largely nationally representative in terms of age, sex and ethnicity, covering around 7% of the UK population. Given the data on demographics, diagnoses, therapies, and tests together with its linkage to other health-related databases, the CPRD is a valuable source of healthcare data [57]. Because CPRD contains detailed personal information, the dataset is not readily available to the public, and its usage depends on approval from the CPRD Research Ethics Committee [58].

For this work, we only considered practices providing healthcare data that met research quality standards within the period from 01/January/1985 to 31/December/2014, and whose records were linked

Table 2

Statistics of CPRD dataset from 01/January/1985 to 31/December/2014. (SD: Standard Deviation, IQR: Interquartile Range, D:Diagnosis, M:Medication).

Number of patients	4, 272, 833
Number of visits	283, 996, 690
Number of visits per patient, Mean (SD)	66.47 (91.1)
Number of visits per patient, Median (IQR)	30 (79)
Number of medical codes	326 (D:222, M:104)
Number of codes in a visit, Mean (SD)	2.34 (1.92)
Number of codes in a visit, Median (IQR)	2 (2)

to the Hospital Episode Statistics database. Furthermore, we focused on patients aged 16 years or older who have been registered with their GP for at least 1 year. This resulted in a dataset of 4, 272, 833 patients, and is profiled in [Table 2](#).

2.3. Modeling pipeline

While CPRD contains many data fields, we limited our analyses to two main scenarios: demographics + diagnoses, and demographics + diagnoses + medications (denoted as DD and DDM, respectively), as these are common predictors that were used by all models in [Table 1](#), and were very likely to exist in all EHR systems. The demographics variables that we took into account were sex (binary), age (continuous), ethnicity (categorical with 12 classes), and region (categorical with 10 classes). In CPRD, diagnoses were coded using Read codes [59] (for primary care) and ICD-10 codes [60] (for Hospital Episode Statistics), while medications were coded using prodcodes (CPRD unique code for the treatment selected by the GP). In their raw format, these coding schemes were unsuitable to work with given their high cardinality, i.e., there were around 110,000 Read codes, 14,000 ICD-10 codes and 55,000 prodcodes. For this reason, we mapped diagnoses from Read and ICD-10 to the CCS (Clinical Classifications Software) [61] coding scheme. For medications, we mapped the prodcodes to level 2 of BNF (British National Formulary) [62] codes. These mappings are known as “clinical groupers” and are commonly used in clinical tasks to reduce the number of medical codes commonly found in EHR.

In our work, we decided to use CCS because it was the most common approach in the papers that we compared. Furthermore, this approach helps address the sparsity issue that is common in medical records, and makes statistical analyses and reporting easier [61]; defining the diseases at a higher granularity will lead to lower frequencies of their occurrences and hence hampers the models’ ability of learning meaningful patterns about them. Furthermore, using clinical groupers can help avoid the memory issues, as raised by Choi et al. [45,50], and from a practical point of view, using the full vocabulary of ICD codes (i.e., more than 10,000 diseases) is not a tractable option for the eNRBM and Deep Patient models, given that such architectures do not use embedding layers and therefore, the memory in the GPUs does not have enough space to allocate such information. In addition to this, diagnoses and medications with a cumulative sum of less than 0.5% of the total number of medical codes (recorded from 01/January/1985 to 31/December/2014), were grouped into the *Drare* and *Mrare* categories, respectively. This resulted in a total of 222 codes for diagnoses and 104 for medications.

For both the DD and DDM scenarios, we focused on two separate outcomes: predicting emergency admission, or heart failure. We chose these two outcomes to capture two important but diverse clinical prediction challenges. Overall, emergency hospital admission is of great importance in healthcare service delivery. Predicting such an outcome helps in planning to ensure services are available to meet unscheduled hospitalisations [63]. The reader can find the emergency admission codes used in this study in [Appendix B](#). Furthermore, heart failure is an important cause of mortality and morbidity, which could be associated

with poor prognosis. Recent trends suggest that the burden of heart failure is increasing, not least because of the trend towards an ageing population [64]. In this study we focused on CCS code 108, i.e. Congestive Heart Failure; Nonhypertensive. The reader can find the related ICD codes in [Appendix C](#). In both cases, we considered a prediction window of 6 months, and the outcome was treated as a binary classification task. Patients were split into separate groups, randomly keeping 53% for training of DL models, 7% for training the top layer classifier in the modular DL architectures (i.e., a random forest in our study), 20% for validation (hyperparameter tuning), and 20% for testing. All groups included data from 01/January/1985, but the time window was different for each of them. The training group covered until 31/December/2012, the validation group until 31/December/2013 and the test group until 31/December/2014. The purpose of this time-splitting is to have a form of external validation (i.e., out of time, out of sample) that is typical in clinical studies, where we want to predict future events based on everything we know up to a certain baseline. In all cases, the prediction interval was used to create the target variable and corresponded to the last six months previous to the end of the corresponding time window (i.e. whether a patient had an emergency admission, or whether he had a heart failure during these six months). Overall, the ratio of emergency admission cases is 0.011, while the ratio of heart failure cases is 0.0015. The statistics of the outcome variables are shown in [Tables 3,4](#). Data before these six-month windows are used for feature generation (see a visual representation of this approach in [Fig. 2](#)). Patients who died before the corresponding prediction interval are excluded from the analysis. For heart failure cases, we did not focus on new incidences since we wanted to identify any heart failure presentation (whether first event or recurrence) given the patient’s history up to a baseline.

While using the same source of data and the same inclusion and exclusion criteria, each model in [Table 1](#) requires a different format for its input data, as described below:

- *eNRBM*: Each patient’s medical history (i.e., the data before dashed line in [Fig. 2](#)) was split into five non-overlapping intervals: (0 – 3), (3 – 6), (6 – 12), (12 – 24), and (24+) months before the prediction window. The information in each interval is formatted as a long sparse vector of length equal to 222 for DD scenario and 326 for DDM scenario, where each entry corresponds to the number of times a single disease (or medication) was diagnosed (or prescribed) in the given interval. All categorical features are converted to dummy variables, and age is scaled between 0 and 1 using information from the training set. This resulted in a sparse vector that, together with demographics information, has 1,135 input variables for DD scenario, and 1,655 for DDM scenario.
- *Deep Patient*: Each patient’s medical history was aggregated into a sparse vector, where each entry corresponds to the number of times a single disease (or medication) was diagnosed (or prescribed). All categorical features are converted to dummy variables, and age is scaled between 0 and 1 using information from the training set. Together with demographic information, this resulted in 247 input variables for DD scenario, and 351 for DDM scenario.
- *DeepR*: It deals with the sequence of events (i.e., diagnoses and

Table 3

Statistics of the outcome variables for the different data partitions for the Demographics + Diagnoses scenario.

Data partition	# of Emergency Admission cases	# of Heart Failure cases
Training (53%)	20,430 (1.08%)	2,754 (0.15%)
Training (modular models) (7%)	2,773 (1.08%)	378 (0.15%)
Validation (20%)	8,262 (1.14%)	1,114 (0.15%)
Testing (20%)	7,115 (0.98%)	982 (0.14%)

Table 4
Statistics of the outcome variables for the different data partitions for the Demographics + Diagnoses + Medications scenario.

Data partition	# of Emergency Admission cases	# of Heart Failure cases
Training (53%)	23,200 (1.18%)	2,764 (0.14%)
Training (modular models) (7%)	3,064 (1.16%)	385 (0.15%)
Validation (20%)	9,308 (1.24%)	1,032 (0.14%)
Testing (20%)	7,998 (1.07%)	947 (0.13%)




Fig. 2. Study Design. A total of 4,272,833 patients were split into training, validation and test sets. A small portion of the training (7%) was kept apart for the modular deep learning architectures. Red region corresponded to the prediction interval, which consisted of the last six months of the respective time window. This region was used to create the two separate outcomes (i.e. whether or not a patient had an emergency admission, or a heart failure during these six months). All data previous to these six months (marked with a dashed line) was used for feature extraction. Patients who died before the corresponding prediction interval are excluded from the analysis. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

medications) directly, while discretising the time between two events as: (0 – 1], (1 – 3], (3 – 6], (6 – 12], and 12+ months, e.g. (0 – 1] corresponds to a duration that goes from 0 (exclusive) to 1 (inclusive) months between medical events. Each time interval is assigned a unique identifier, which is treated as a word; for instance, 0–1m is a word for the (0 – 1] interval gap. Each patient’s medical history is expressed as a long sentence of medical codes, where the duration between two consecutive visits will appear in the sequence as one of the five preserved interval words. Similar to the approach of Nguyen et al. [20], the sentences were trimmed to keep the last 100 words in order to avoid the effects of some patients, who have very long sentences, that could severely skew the data distribution.

- **RETAIN:** Each patient’s information was split in three components. The first one was made of a list of sublists, where each sublist contained all the medical codes that were recorded on a single date. These sublists were sequentially ordered from oldest to newest. The second component consisted of a list with the demographics information. Finally, the third component, was made of a list of values indicating the difference in days between consecutive visits. These values were ordered from oldest to newest, and the first of these differences was set to zero, given that this information was not available for the first visit.

This input formatting was carried out after the data partitioning process. A short description of the inclusion criteria and outcomes of interest of the selected works are shown in the Appendix D. For more

details, readers are directed to the corresponding publications.

All four models selected were coded based on the description found in the corresponding publications, except for the RETAIN model, which has the code available online [65]. Note that the original model in Nguyen et al. [20] for the Deepr architecture only considered dynamic information as the input of the convolutional layer (after embedding). However, for the sake of a fair comparison with the other selected models, we had to include static information (i.e. demographics) into the network. This resulted in a slight modification of the original architecture, via the parallel processing of two kinds of inputs (multi-input network). This adjustment amounted to merging two networks: the first (processing dynamic information) had the same structure as Deepr, whereas the second (processing static information) consisted of one fully connected layer. The two intermediate outputs were then merged together through a further fully connected layer, before the final classification task (see Fig. 3).

The implementations were done with PyTorch 0.4.1 [66], Keras 2.2.2 [67] and scikit-learn 0.20.0 [68] using two NVIDIA Titan Xp graphics cards. Furthermore, all DL models selected have several hyperparameters as shown in Appendix E. In order to tune these hyperparameters, we used a tree-based Bayesian Optimisation approach [69,70], as implemented in Scikit-Optimize [71]. We chose a batch size of 1,024 patients for all models, except for RETAIN, where we used a batch size of 256 patients due to memory constraints in the GPUs. To benchmark our results, we used bag-of-words representation with logistic regression (BOW + LR), and a simple RBM with the same input format as in Tran et al. [40]. For the random forest in the modular DL models, we used 200 trees while the remaining hyperparameters were optimised using the same tree-based Bayesian Optimisation approach¹.

2.4. Performance metrics

Model performance was evaluated using the area under the ROC curve (AUROC), similar to what the majority of the articles in Table 2 reported. The AUROC is a discrimination metric equivalent to the probability that a classification model will rank a randomly chosen positive instance higher than a randomly chosen negative instance [72,73]. Thus, the AUROC ranges from 0.5 (no discrimination ability) to 1 (perfect discrimination). However, this metric can be misleading when the classes are very imbalanced [74–76]. To avoid this, the area under the precision-recall curve (AUPRC) is useful as it considers the tradeoff between precision (a measure of result relevancy) and recall (a measure of how many truly relevant results are returned) for different thresholds [77,78]. The AUPRC range goes from 0 to 1, with higher values preferred overall [79]. In addition, the F1-Score is also taken into account given that it is a weighted average of the precision and recall, ranging from 0 (worst score) to 1 (best score) [80].

3. Results

The performances of all the tested models are shown in Tables 5,6 for the DD scenario, and Tables 7,8 for the DDM scenario. The models were trained 5 times independently to provide summary statistics and 95% confidence intervals. In general, the RETAIN model has the best performance when compared to other DL models; in both scenarios, and across a range of metrics. This agrees with the results presented in Choi et al. [45] in terms of AUROC, and emphasises the advantage of using RNN to exploit the sequential nature of EHR. In addition, the inclusion of medications provides an improvement in the AUROC and AUPRC for the majority of the models in these two prediction tasks. The best hyperparameters found for all the DL and baseline models are shown in Appendix E.

When compared to simpler techniques such as BOW + LR, all DL

¹ For the interested readers, our codes will be shared upon request.




Fig. 3. Modified DeepPr architecture in order to consider demographics and have a fair comparison with the other selected models. EHR are processed through a series of steps that include sequencing, embedding, convolution, pooling and classification. Static information (i.e., demographics) is merged together and passed through a fully connected layer before the final classification task.

models show better performance. For instance, when considering Emergency Admission in the DD scenario, the RETAIN model outperforms the BOW + LR model by 25% in AUROC, 264% in AUPRC, and 174% in F1-Score. Similarly, for Heart Failure cases, the RETAIN model outperforms the same by 18% in AUROC, 490% in AUPRC, and 350% in F1-Score. On the other hand, when considering Emergency Admission in the DDM scenario, the RETAIN model outperforms the BOW + LR model by 31% in AUROC, 336% in AUPRC, and 183% in F1-Score. Similarly, for Heart Failure cases, the RETAIN model outperforms the same by 39% in AUROC, 800% in AUPRC, and 515% in F1-Score. Overall, this is likely due to the ability of the DL models to reduce a complex input space (made out of hundreds of diagnosis and medication codes) to a small yet predictive representation of the EHR sequence for predicting either emergency admission or heart failure. Similarly, when compared with the simple RBM model, the eNRBM has a similar performance in both classification tasks, as expected, given the similar layout of both models (see Table E.11). Further, both Deep Patient and DeepPr have a slightly worse performance compared to the simple RBM when predicting emergency admission, but this is improved when predicting heart failure. Finally, the RETAIN model has the best performance in both tasks for the DD and DDM scenarios.

4. Discussions and future work

DL is becoming the ubiquitous approach for the analysis of EHR data, because of the models' ability to process huge amounts of data without the need to perform explicit feature engineering by domain

Table 5
Comparison for the Demographics + Diagnoses scenario (Emergency Admission).

Model	AUROC	AUPRC	F1-Score
eNRBM	0.803 (0.803–0.804)	0.051 (0.051–0.051)	0.052 (0.052–0.052)
Deep Patient	0.801 (0.801–0.802)	0.047 (0.047–0.047)	0.053 (0.053–0.053)
DeepPr	0.815 (0.813–0.817)	0.053 (0.052–0.054)	0.112 (0.107–0.117)
RETAIN	0.822 (0.819–0.826)	0.062 (0.061–0.064)	0.118 (0.109–0.127)
BOW + LR	0.654 (0.627–0.681)	0.017 (0.015–0.018)	0.043 (0.031–0.054)
RBM	0.807 (0.807–0.808)	0.050 (0.050–0.050)	0.053 (0.053–0.053)

* Data represented as: Mean (95% Confidence Interval).

Table 6
Comparison for the Demographics + Diagnoses scenario (Heart Failure).

Model	AUROC	AUPRC	F1-Score
eNRBM	0.912 (0.911–0.913)	0.018 (0.017–0.018)	0.017 (0.017–0.017)
Deep Patient	0.948 (0.948–0.949)	0.040 (0.039–0.041)	0.024 (0.024–0.024)
DeepPr	0.938 (0.907–0.969)	0.049 (0.036–0.063)	0.110 (0.089–0.131)
RETAIN	0.951 (0.949–0.952)	0.065 (0.061–0.069)	0.135 (0.126–0.144)
BOW + LR	0.801 (0.775–0.827)	0.011 (0.009–0.012)	0.030 (0.025–0.034)
RBM	0.897 (0.897–0.898)	0.016 (0.016–0.017)	0.015 (0.015–0.015)

* Data represented as: Mean (95% Confidence Interval).

Table 7
Comparison for the Demographics + Diagnoses + Medications scenario (Emergency Admission).

Model	AUROC	AUPRC	F1-Score
eNRBM	0.831 (0.831–0.832)	0.071 (0.071–0.071)	0.063 (0.062–0.063)
Deep Patient	0.813 (0.813–0.813)	0.060 (0.060–0.061)	0.059 (0.059–0.059)
DeepPr	0.829 (0.828–0.831)	0.069 (0.067–0.071)	0.131 (0.118–0.144)
RETAIN	0.847 (0.845–0.849)	0.083 (0.082–0.083)	0.153 (0.151–0.154)
BOW + LR	0.646 (0.576–0.717)	0.019 (0.015–0.023)	0.054 (0.046–0.063)
RBM	0.840 (0.840–0.840)	0.072 (0.072–0.073)	0.066 (0.066–0.066)

*Data represented as: Mean (95% Confidence Interval).

Table 8
Comparison for the Demographics + Diagnoses + Medications scenario (Heart Failure).

Model	AUROC	AUPRC	F1-Score
eNRBM	0.920 (0.920–0.921)	0.020 (0.019–0.021)	0.014 (0.014–0.014)
Deep Patient	0.947 (0.947–0.948)	0.040 (0.039–0.041)	0.023 (0.022–0.023)
DeepR	0.949 (0.947–0.952)	0.039 (0.032–0.046)	0.085 (0.049–0.120)
RETAIN	0.950 (0.946–0.954)	0.054 (0.053–0.056)	0.117 (0.098–0.136)
BOW + LR	0.682 (0.613–0.752)	0.006 (0.002–0.009)	0.019 (0.011–0.027)
RBM	0.917 (0.917–0.917)	0.023 (0.022–0.023)	0.014 (0.014–0.014)

* Data represented as: Mean (95% Confidence Interval).

experts unlike more traditional statistical models. In this work, we carried out a comparative review of the common DL architectures for EHR, to compare their properties while standardising the common sources of paper-to-paper variability (e.g., dataset, sample size, medical codes, clinical questions, and performance metrics). To the best of our knowledge, this is the first paper that has carried out such research.

Most of the DL research in this field, however, analysed EHR data that came from private hospitals [40,38,20,45]. These datasets vary in terms of sample size and the information they contain (demographics, clinical data, vital signs, laboratory tests and medications, among others), and hence DL's performance on them might not be reproduced across other EHR datasets (and hence different populations). Among EHR data sources, CPRD is one of the largest accessible EHR datasets for a patient population that is largely representative of the UK population in terms of age, sex, and ethnicity. The CPRD not only includes data on demographics, diagnoses, therapies, and tests, but it is also linked to data on hospitalisations, outpatient visits, accident and emergency attendances, diagnostic imaging, death registration, cancer registration, and socioeconomic status [57]. It is a valuable source of healthcare data that enables models to be trained in a real healthcare ecosystem.

Although CPRD has been around for over a decade with hundreds of publications and widespread use in epidemiology [81–83,64,84], its potential for machine learning and DL is currently untapped and its use only in early stages (i.e., Rahimian et al. [19]). With this study being the first to employ DL on CPRD, we aimed to show the richness of this database for various types of DL frameworks, given its large sample size, variety of information, and representativeness of the UK population, which can be used on a large range of analyses on risk predictions and disease trajectories.

In this work, we used four models that covered the main DL architectures found in the literature. The eNRBM and Deep Patient models were trained in a modular way. First, they aimed to get a meaningful unsupervised representation that identifies hierarchical regularities and dependencies in the input data; such a representation is then used as input for a supervised learning algorithm (i.e., a random forest in our study). Given both models' inability to deal with sequential data explicitly, they require to represent a patient's journey as a sparse vector, which ignores the elapsed time between visits. In practice, however, time between and since events is clinically relevant. For instance, Rahimian et al. [19] have shown a 10.8% higher AUC (0.848 with a gradient boosting classifier compared to 0.740 with a Cox proportional hazards model) for prediction of risk of emergency admission within 24 months using engineering features that represented time since diagnoses. Thus, most recent developments in the use of DL for modelling EHR data has been using sequential models such as RNN.

DeepR and RETAIN models, on the other hand, are trained using an end-to-end approach, directly connecting the input (e.g., medical codes and other information) to the outcome of interest. This can be an advantage if the aim is to identify meaningful patterns for a specific clinical question, but if this is not the case, a different model needs to be trained for each given outcome. Nevertheless, since these models incorporate the elapsed time between visits, they are more suitable for modelling patient trajectories. While DeepR employed an interesting

approach to characterise time between events by defining special words that represent time intervals, RETAIN considered time between visits as an additional feature. Overall, the RETAIN model has the best performance in both tasks for the DD and DDM scenarios. This highlights the strength of RNN and attention mechanisms for analysing sequential data, thus various modifications to RNN are currently the state of the art research in DL for EHR.

Regardless of which classification model one uses, one of the most common ways to assess such model performance is the AUROC. In our study, we achieved a high AUROC that is comparable to what other papers reported [20,45,48]. Furthermore, we employed an additional metric, AUPRC, which most papers in the literature did not report, and with respect to which our models performed poorly. Such observations have been previously reported by others, e.g., Davis and Goadrich [85], when datasets were highly imbalanced. It is possible that focusing on more specific or relevant clinical subgroups may improve the prediction as models are likely to incorporate more relevant parameters for the condition of interest than for all the patient population. This highlights the difficulty that even these models encounter when dealing with imbalanced clinical tasks. Such a scenario is typical in many healthcare applications where the minority class is buried by the majority class. Several approaches have been proposed for dealing with imbalanced datasets, including the work by Chawla et al. [86] and Lemaitre et al. [87], but further research is required, particularly for DL models.

One of the key strengths of DL models is their ability to map one's full medical record to a low-dimensional representation. In our analyses, the patient representations (projected onto 2D using a UMAP (Uniform Manifold Approximation and Projection) [88] for visual clarity) are shown in Figs. 4 and 5. Note that the representations from modular models tend to form clusters, while the representations from the end-to-end models seem to form a continuum. Nevertheless, it can be observed that cases of emergency admission (in blue) and heart failure (in red) tend to get grouped in separable regions. Further profiling of these clusters can help identify and assess common characteristics of patients to improve the performance of the models and to guide future clinical research.

As part of the objectives of this work to provide a guideline to work with DL for EHR, we identified that the main issues are related to the choice of inclusion criteria and preprocessing of the data. The inclusion criteria usually depends on the clinical question and outcome of interest; it is a key part of a good study design. Furthermore, it is of great importance to properly split the data into the usual training, validation and test sets as this would help to avoid overfitting and properly evaluate the DL models. In terms of data preprocessing, we highlighted the importance of using clinical groupers to reduce the cardinality of both diagnoses and medications. This could be thought as a form of dimensionality reduction that prevents very long and sparse input vectors that would hamper the training process of the DL models. In this work, we used the CCS coding scheme, but other clinical [89] and drug [90] classification systems could be used depending on the study requirements.

One important consideration when working with DL models is the tuning of the corresponding hyperparameters (as shown in Appendix

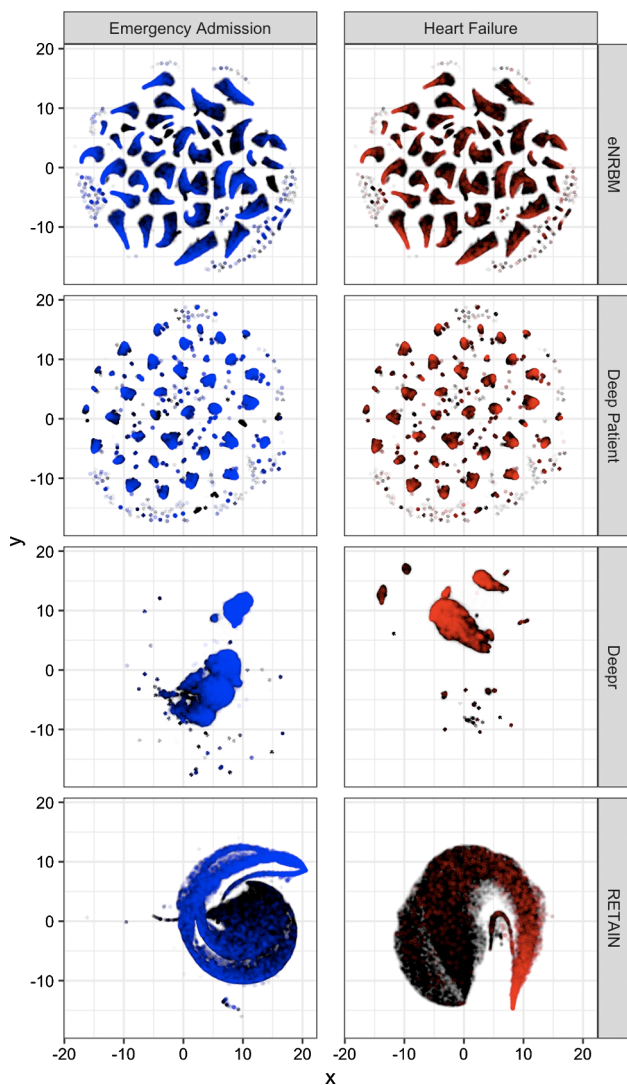


Fig. 4. Patient representations obtained by each DL model in the DD scenario, after projection to 2D using UMAP. Each point corresponds to a patient in the test set; blue and red dots correspond to cases of emergency admission or heart failure, respectively. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

E). In this work, we employed a tree-based Bayesian Optimisation approach [69,70], as implemented in Scikit-Optimize [71], to find the most appropriate values. This approach takes fewer evaluations to achieve a similar result compared with exhaustive grid search and random search given that it learns from its evaluation history and suggests better hyperparameters to test as the search progresses. Still, Bayesian optimisation can be time-consuming as the models become larger and more complex, so a lot of research has been done to make this process more robust and scalable [91–94]. Nowadays, new best practices are proposed constantly in this fast-evolving field, so it is hard to provide a general set of rules for efficient hyperparameter tuning. The work by Howard et al. [95] has been highly influential as the authors suggest that the learning rate is the most important hyperparameter when training a DL model, and advocate for a more iterative approach using techniques like differential learning rates, cyclical learning rates [96], learning rate schedulers, and stochastic gradient descent with restarts [97], which seem to work irrespectively of the DL architecture.

Overall, there is no DL model that dominates the current state-of-the-art for the analysis of EHR, and that outperforms all expert-driven models yet (as shown in the works of Jacobs et al. [98] and Rahimian

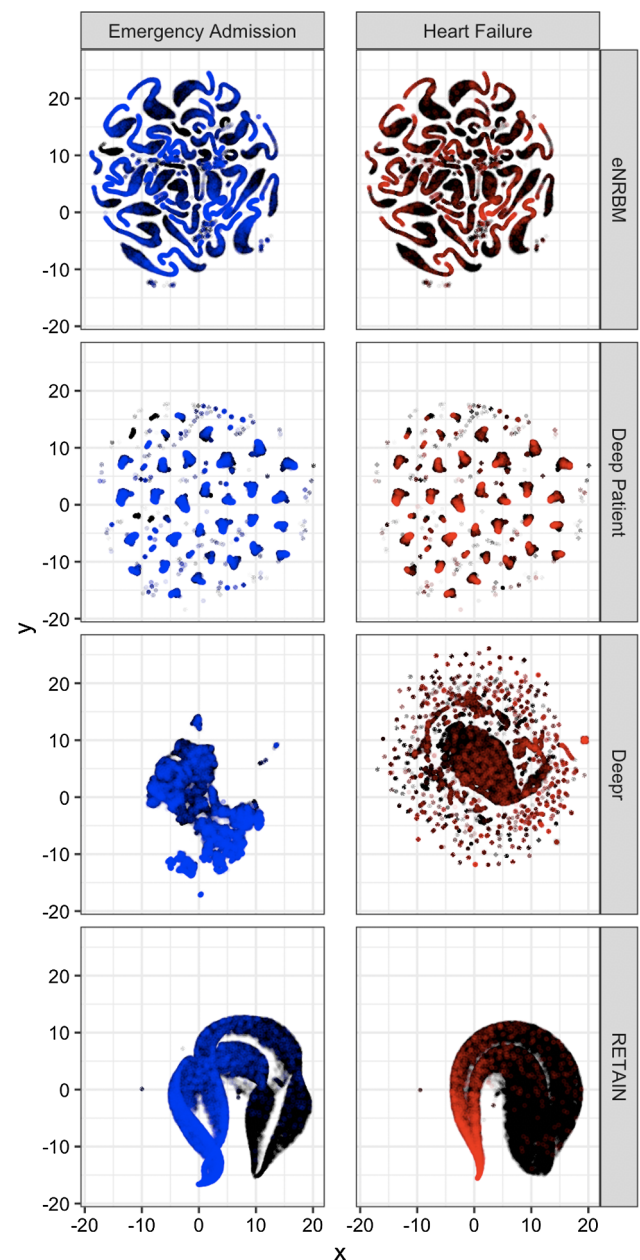


Fig. 5. Patient representations obtained by each DL model in the DDM scenario, after projection to 2D using UMAP. Each point corresponds to a patient in the test set; blue and red dots correspond to cases of emergency admission or heart failure, respectively. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

et al. [19]). Furthermore, DL models remain difficult to interpret and unable to provide uncertainty estimation, which is undesirable in clinical settings. Although, the works by Gal et al. [99,100], Tran et al. [101] and Chakraborty et al. [102] have contributed to overcome these issues, further research is required.

Finally, there seems to be a trend towards the use of RNN architectures given the sequential nature of medical records. Several works have been developed recently where RNN are applied to a series of case studies like early detection of sepsis [103], hospital readmission for lupus patients [104], and detection of adverse medical events in EHR [105]. Furthermore, the field of natural language processing has greatly advanced as evidenced by the works of Transformers [106,107], contextualised word representations [108], and transfer learning for language models [109]. Given the similarities between natural language

processing tasks and EHR, the application of these techniques provides interesting research directions that could result in the development of medical embeddings for analysis of EHR [110,111], understanding of disease clusters and trajectories [112], generation of synthetic EHR [113], and the use of pretrained models to avoid the need for starting from scratch every time a new clinical question comes. The application of DL to EHR seems to have an exciting future.

5. Conclusions

The use of DL to analyse EHR data has increased over the past years; a growth that is continuing to be facilitated by the availability of more data (EHR and beyond), developments in DL (specifically, models for sequential data), and innovative ways of combining these two trends. In this work, we implemented key deep learning architectures to learn an efficient patient representation for predicting emergency admission, and heart failure. Our objective here was to help the field have a comparative view of these approaches, and to assess their strengths and weaknesses when it comes to EHR.

Along this work, we introduced CPRD, which is one of the world's largest primary care databases, and showed how data from primary care can provide predictions that can be of value in policy and practice of care. Given the complexity of primary care EHR (heterogeneous events recorded in irregular intervals with varying degree of richness and quality across different individuals), and its importance in provision of care in many parts of the world, we believe that the set of best practices we shared for them (e.g., inclusion criteria, preprocessing, medical codes/grouping, performance metrics, and hyperparameter tuning) will be of great value in helping DL research in EHR.

Our work showed the strength of recurrent neural networks in dealing with the temporal nature of the EHR. This was consistent with the developments in modelling EHR-like data from other domains, such as natural language and time series data. Our future research aims to explore techniques and methodologies from such domains, and apply them to other types of data from different healthcare systems.

Appendix A. Abbreviations

The following abbreviations are used throughout the document:

- EHR: Electronic Health Records
- SVM: Support Vector Machines
- DL: Deep Learning
- FFNN: Feed-Forward Neural Networks
- CNN: Convolutional Neural Networks
- RNN: Recurrent Neural Networks
- eNRBM: Electronic Medical Records-driven nonnegative restricted Boltzmann machines
- SDA: Stacked Denoising Autoencoder
- RETAIN: Reverse Time Attention
- CPRD: Clinical Practice Research Datalink
- GP: General Practice
- ICD: International Statistical Classification of Diseases and Related Health Problems
- ACHI: Australian Classification of Health Interventions
- ACS: Australian Coding Standard
- ATC: Anatomical Therapeutic Chemical Classification
- CPT: Current Procedural Terminology
- NDC: National Drug Codes
- GPI: Generic Product Identifier
- KCD: Korean Statistical Classification of Diseases and Related Health Problems
- SD: Standard Deviation
- IQR: Interquartile Range
- BOW: Bag of Words
- LR: Logistic Regression
- AUROC: Area Under the Receiver Operating Characteristic Curve
- AUPRC: Area Under the Precision-Recall Curve
- ND: Not Defined

6. Author contributions statement

G.S.K. conceived the idea for this work; J.R.A.S worked on the data cleaning and formatting, as well as on the eNRBM, Deep Patient and RETAIN models; F.E.D.R worked on the Deepr model; Y.Z. reviewed the data pipelines and model pipelines for eNRBM, Deep Patient, Deepr, and RETAIN; J.R.A.S reported the results for the eNRBM, Deep Patient and RETAIN models; F.E.D.R reported the results for the Deepr model; J.R.A.S wrote the rest of the paper; G.S.K. and Y.Z. contributed to the discussion part of the paper. All authors contributed to multiple parts of the manuscript review, as well as the style and overall contents.

Declaration of Competing Interest

None.

Acknowledgements

This research was funded by the Oxford Martin School (OMS) and supported by the National Institute for Health Research (NIHR) Oxford Biomedical Research Centre (BRC). It also got support from the PEAK Urban programme, funded by UK Research and Innovation's Global Challenge Research Fund, Grant Ref: ES/P011055/1. The views expressed are those of the authors and not necessarily those of the OMS, the UK National Health Service (NHS), the NIHR or the Department of Health and Social Care.

This work used data provided by patients and collected by the NHS as part of their care and support and would not have been possible without access to this data. The NIHR recognises and values the role of patient data, securely accessed and stored, both in underpinning and leading to improvements in research and care.

We also thank Wayne Dorrington for his help in making the graphical abstract that summarises the contents of this work, and Emma Copland for revising the first draft of this manuscript.

- DD: Demographics + Diagnoses
- DDM: Demographics + Diagnoses + Medications

Appendix B. Emergency Admission Codes

Table B.9 shows the Medcodes considered to identify emergency admissions.

Table B.9

Medcodes considered to identify emergency admissions.

Medcode	Description
140	Cauterisation of internal nose
314	Diagnostic endoscopic examination and biopsy lesion larynx
1081	Herpes simplex disciform keratitis
1158	Other operation on cornea NOS
3488	Rebound mood swings
6269	Herpesviral infection of perianal skin and rectum
6885	Congenital aortic valve stenosis
7058	Admit diabetic emergency
7059	[V]Palliative care
7242	FESS/Post operative division of adhesions
7503	[D]Heart murmur, undiagnosed
8082	Other specified other repair of vagina
8265	Closed fracture lumbar vertebra, wedge
9409	[X]Cut by glass
11413	Other finger injuries NOS
11963	Repair of chest wall
12038	[SO]Leg region
12243	[SO]Tricuspid valve
13706	Stroke/transient ischaemic attack referral
22296	Trucut transperineal biopsy of prostate
22374	Fracture of carpal bone
23106	Serum alkaline phosphatase NO
25057	Seen by ear, nose and throat surgeon
29190	Sensorineural hear loss, unilateral unrestricted hear/contralat side
29988	Gittes endoscopic bladder neck suspension
30027	Fracture of malar and maxillary bones
32898	Peritonitis - tuberculous
35328	Pressure ulcer assessment
37543	Ecstasy poisoning
38379	Mechanical complication of tendon graft
43828	Removal of onychophosis
46824	Congenital malformation of aortic and mitral valves unspecified
48729	Acute duodenal ulcer with haemorrhage and perforation
66764	Closed skull vlt no intracranial injury
67786	Encephalitis due to meningococcus
95163	Labyrinthine round window fistula
99761	Acute lower respiratory tract infection
102555	Tiabendazole poisoning

Appendix C. Heart Failure Codes

Table C.10 shows the ICD codes that are mapped to CCS code 108 (Congestive Heart Failure; Nonhypertensive).

Table C.10

ICD codes mapped to CCS code 108 (Congestive Heart Failure; Nonhypertensive).

ICD	CCS	Description
I0981	108	Rheumatic heart failure
I501	108	Left ventricular failure, unspecified
I5020	108	Unspecified systolic (congestive) heart failure
I5021	108	Acute systolic (congestive) heart failure
I5022	108	Chronic systolic (congestive) heart failure
I5023	108	Acute on chronic systolic (congestive) heart failure
I5030	108	Unspecified diastolic (congestive) heart failure
I5031	108	Acute diastolic (congestive) heart failure
I5032	108	Chronic diastolic (congestive) heart failure
I5033	108	Acute on chronic diastolic (congestive) heart failure
I5040	108	Unspecified combined systolic and diastolic (congestive) heart failure
I5041	108	Acute combined systolic and diastolic (congestive) heart failure
I5042	108	Chronic combined systolic and diastolic heart failure

(continued on next page)

Table C.10 (continued)

ICD	CCS	Description
I5043	108	Acute on chronic combined systolic and diastolic heart failure
I50810	108	Right heart failure, unspecified
I50811	108	Acute right heart failure
I50812	108	Chronic right heart failure
I50813	108	Acute on chronic right heart failure
I50814	108	Right heart failure due to left heart failure
I5082	108	Biventricular heart failure
I5083	108	High output heart failure
I5084	108	End stage heart failure
I5089	108	Other heart failure
I509	108	Heart failure, unspecified

Appendix D. Data: inclusion criteria and outcome

D.1. eNRBM

Tran et al. [40] used a mental health cohort extracted from a large regional hospital in Australia and collected between January 2009 and March 2012. Any patient who had at least one encounter with the hospital services, as well as one risk assessment for suicide, was included. Each assessment was considered as a data point from which a prediction would be made. The sample size was 7, 578 patients (49.3% male, 48.7% under 35), for a total of 17, 566 assessments. The future outcomes within 3 months following an assessment were stratified in 3 ordinal levels of risk: no-risk, moderate-risk (non-fatal consequence), and high-risk (fatal consequence). The risk classes related to a diagnosis were decided in relation to the ICD-10 codes. If there were more than one outcome class, the highest risk class was chosen. There were 86.9% no-risk outcomes, 8.2% moderate-risk and 4.9% high-risk. Since the completed suicides were rare, the class distributions were quite imbalanced.

D.2. Deep Patient

In this work, the authors used the Mount Sinai data warehouse. All patients with at least one diagnosed disease (expressed as numerical ICD-9) between 1980 and 2014, inclusive, were considered in Miotto et al. [38]. This led to a dataset of about 1.2 million patients, with every patient having an average of 88.9 records. All patients with at least one recorded ICD-9 code were split in three independent datasets for evaluation purposes (i.e., every patient appeared in only one dataset). First, the authors retained 81, 214 patients having at least one new ICD-9 diagnosis assigned in 2014 and at least ten records before that. These patients composed validation (i.e., 5, 000 patients) and test (i.e., 76, 214 patients) sets for the supervised evaluation (i.e., future disease prediction). In particular, all the diagnoses in 2014 were used to evaluate the predictions computed using the patient data recorded before December 31, 2013. The requirement of having at least ten records per patient was set to ensure that each test case had some minimum of clinical history that could lead to reasonable predictions. A random subset of 200, 000 different patients was sampled with at least five records before the split-point (December 31, 2013) to use as training set for the disease prediction experiment. Finally, the authors created the training set for the unsupervised feature learning algorithms using the remaining patients having at least five records by December 2013. This led to a dataset composed of 704, 587 patients and 60, 238 clinical descriptors.

D.3. DeepR

Data in Nguyen et al. [20] were collected from a large private hospital chain in Australia from July 2011 to December 2015, using Australian Coding Standards (ICD-10-AM for diagnoses and the Australian Classification of Health Interventions for procedures). Data consisted of 590, 546 records (around 300, 000 patients), each corresponding to an admission (characterised by an admission time and a discharge time). The outcome consisted in unplanned readmission (corresponding to nearly 2% of the sample). The risk group was made of patients with at least one unplanned readmission within 6 months (4, 993 patients), or within 3 months (3, 788 patients) from a discharge, regardless of the admitting diagnosis. For each risk case, a control case was randomly picked from the remaining patients. For each risk/control group, 16.7% of patients were used for model tuning, 16.7% for testing, and the rest for training. A discharge (except for the last one in the risk groups) was randomly selected as the prediction point, from which the future risk would be predicted.

D.4. RETAIN

In Choi et al. [45], the authors focused on heart failure prediction. The dataset consisted of EHR from Sutter Health, including patients from 50 to 80 years of age. Diagnosis, medication and procedure codes were extracted from the encounter records, medication orders, procedure orders and problem lists, and then aggregated into existing medical groupers so as to reduce the dimensionality while preserving the clinical information of the input variables. From the source dataset, 3, 884 cases were selected and approximately 10 controls were considered for each case correspondingly (28, 903 controls). Medical codes were extracted in the 18-month window before the baseline. The patient cohort was divided into the training, validation and test sets in a 0.75: 0.1: 0.15 ratio. The validation set was used to determine the values of the hyperparameters.

Appendix E. Hyperparameters

Tables E.11, E.14, E.16 show the different hyperparameters for each of the selected DL models together with the values identified using a tree-based Bayesian Optimisation approach [69,70], as implemented in Scikit-Optimize [71]. Tables E.12 and E.13 show the corresponding

Table E.11
Hyperparameters' values for modular deep learning architectures.

Model	Hyperparameter	Scenario	
		DD	DDM
eNRBM	# hidden units	234	259
	Learning rate	0.332	0.154
	Nonnegativity cost	0.002	0.001
	Smoothness cost	0.014	0.018
	Epochs	38	20
Deep Patient	# hidden units	181	159
	# hidden layers	1	1
	Learning rate	0.018	0.014
	Noise level	0.004	0.13
	Epochs	14	27
RBM	# hidden units	214	154
	Learning rate	0.451	0.332
	Epochs	25	19

Table E.12
Hyperparameters' values for modular deep learning architectures with emergency admission as the outcome.

Model	Hyperparameter	Scenario	
		DD	DDM
eNRBM	Fraction of features to consider	0.209	0.162
	Minimum number of samples at a leaf node	1206	1363
Deep Patient	Fraction of features to consider	0.185	0.477
	Minimum number of samples at a leaf node	1136	1050
RBM	Fraction of features to consider	0.367	0.440
	Minimum number of samples at a leaf node	1225	1244

Table E.13
Hyperparameters' values for modular deep learning architectures with heart failure as the outcome.

Model	Hyperparameter	Scenario	
		DD	DDM
eNRBM	Fraction of features to consider	0.416	0.397
	Minimum number of samples at a leaf node	1153	4976
Deep Patient	Fraction of features to consider	0.493	0.396
	Minimum number of samples at a leaf node	1657	1285
RBM	Fraction of features to consider	0.467	0.286
	Minimum number of samples at a leaf node	2170	3605

Table E.14
Hyperparameters' values for end-to-end deep learning architectures with emergency admission as the outcome.

Model	Hyperparameter	Scenario	
		DD	DDM
DeepR	Filters	45	31
	Kernel size	5	3
	Fully connected units - dynamic info	25	34
	Fully connected units - static info	41	46
	Fully connected units - after merging	8	6
	Dropout - dynamic info	0.256	0.423
	Dropout - static info	0.112	0.117
	Dropout - after merging	0.114	0.439
	Positive class weight	53.022	78.120
	Learning rate	0.002	0.001

(continued on next page)

Table E.14 (continued)

Model	Hyperparameter	Scenario	
		DD	DDM
RETAIN	Embedding size	135	150
	Recurrent layer size	113	79
	Dropout index	0.333	0.482
	Dropout context	0.302	0.104
	L2 regularization	0.048	0.012
	Epochs	19	8

Table E.15

Hyperparameters' values for baseline models with emergency admission as the outcome.

Model	Hyperparameter	Scenario	
		DD	DDM
BOW + LR	ℓ_1 Regularization coefficient	1.77e-08	1.12e-06
	Positive class weight	60	34

Table E.16

Hyperparameters' values for end-to-end deep learning architectures with heart failure risk as the outcome.

Model	Hyperparameter	Scenario	
		DD	DDM
DeepPr	Filters	45	43
	Kernel size	4	4
	Fully connected units - dynamic info	9	50
	Fully connected units - static info	37	24
	Fully connected units - after merging	26	25
	Dropout - dynamic info	0.244	0.406
	Dropout - static info	0.185	0.371
	Dropout - after merging	0.349	0.299
	Positive class weight	67.151	1.969
	Learning rate	0.004	0.005
RETAIN	Embedding size	193	143
	Recurrent layer size	104	116
	Dropout index	0.500	0.457
	Dropout context	0.043	0.047
	L2 regularization	0.008	0.009
	Epochs	11	14

Table E.17

Hyperparameters' values for baseline models with heart failure as the outcome.

Model	Hyperparameter	Scenario	
		DD	DDM
BOW + LR	ℓ_1 Regularization coefficient	1.99e-06	8.97e-03
	Positive class weight	97	73

hyperparameters for the random forest used as the top layer classifier in the modular DL architectures. Tables E.15 and E.17 show the corresponding hyperparameters and the identified values for the baseline models.

References

- [1] G.S. Birkhead, M. Klompas, N.R. Shah, Uses of electronic health records for public health surveillance to advance public health, *Annu. Rev. Public Health* 36 (1) (2015) 345–359, <https://doi.org/10.1146/annurev-publhealth-031914-122747>.
- [2] L.L. Weed, Medical records that guide and teach, *New Engl. J. Med.* 278 (1968) 652–657.
- [3] T. Botsis, G. Hartvigsen, F. Chen, C. Weng, Secondary use of EHR: data quality issues and informatics opportunities, *Summit Transl. Bioinform.* 2010 (2010) 1.
- [4] P.B. Jensen, L.J. Jensen, S. Brunak, Mining electronic health records: towards better research applications and clinical care, *Nat. Rev. Genet.* 13 (6) (2012) 395.
- [5] B. Shickel, P.J. Tighe, A. Bihorac, P. Rashidi, Deep EHR: a survey of recent advances in deep learning techniques for Electronic Health Record (EHR) analysis, *IEEE J. Biomed. Health Inform.* 22 (5) (2018) 1589–1604, <https://doi.org/10.1109/JBHI.2017.2767063>.
- [6] S.M. Meystre, G.K. Savova, K.C. Kipper-Schuler, J.F. Hurdle, Extracting information from textual documents in the electronic health record: a review of recent research, *Yearbook Med. Inform.* 17 (01) (2008) 128–144.
- [7] M. Jiang, Y. Chen, M. Liu, S.T. Rosenbloom, S. Mani, J.C. Denny, H. Xu, A study of machine-learning-based approaches to extract clinical entities and their assertions from discharge summaries, *J. Am. Med. Inform. Assoc.* 18 (5) (2011) 601–606.

- [8] F. Doshi-Velez, Y. Ge, I. Kohane, Comorbidity clusters in autism spectrum disorders: an electronic health record time-series analysis, *Pediatrics* 133 (1) (2014) e54–e63, <https://doi.org/10.1542/peds.2013-0819>.
- [9] L. Li, W.-Y. Cheng, B.S. Glicksberg, O. Gottesman, R. Tamler, R. Chen, E.P. Bottinger, J.T. Dudley, Identification of type 2 diabetes subgroups through topological analysis of patient similarity, *Sci. Transl. Med.* 7 (311) (2015) 311ra174.
- [10] S. Ebadollahi, J. Sun, D. Gotz, J. Hu, D. Sow, C. Neti, Predicting patient's trajectory of physiological data using temporal trends in similar patients: a system for near-term prognostics, in: *AMIA Annual Symposium Proceedings*, Vol. 2010, American Medical Informatics Association, 2010, p. 192.
- [11] D. Zhao, C. Weng, Combining PubMed knowledge and EHR data to develop a weighted bayesian network for pancreatic cancer prediction, *J. Biomed. Inform.* 44 (5) (2011) 859–868.
- [12] P.C. Austin, J.V. Tu, J.E. Ho, D. Levy, D.S. Lee, Using methods from the data-mining and machine-learning literature for disease classification and prediction: a case study examining classification of heart failure subtypes, *J. Clin. Epidemiol.* 66 (4) (2013) 398–407.
- [13] G.J. Kuperman, A. Bobb, T.H. Payne, A.J. Avery, T.K. Gandhi, G. Burns, D.C. Classes, D.W. Bates, Medication-related clinical decision support in computerized provider order entry systems: a review, *J. Am. Med. Inform. Assoc.* 14 (1) (2007) 29–40.
- [14] R. Miotto, C. Weng, Case-based reasoning using electronic health records efficiently identifies eligible patients for clinical trials, *J. Am. Med. Inform. Assoc.* 22 (e1) (2015) e141–e150.
- [15] F.E. Harrell Jr., K.L. Lee, R.M. Califf, D.B. Pryor, R.A. Rosati, Regression modelling strategies for improved prognostic prediction, *Stat. Med.* 3 (2) (1984) 143–152.
- [16] I. Kurt, M. Ture, A.T. Kurum, Comparing performances of logistic regression, classification and regression tree, and neural networks for predicting coronary artery disease, *Exp. Syst. Appl.* 34 (1) (2008) 366–374.
- [17] R.J. Carroll, A.E. Eyler, J.C. Denny, Naïve electronic health record phenotype identification for rheumatoid arthritis, in: *AMIA Annual Symposium Proceedings*, Vol. 2011, American Medical Informatics Association, 2011, p. 189.
- [18] J. Hippisley-Cox, C. Coupland, Predicting risk of emergency admission to hospital using primary care data: derivation and validation of QAdmissions score, *BMJ Open*, vol. 3, 8.
- [19] F. Rahimian, G. Salimi-Khorshidi, J. Tran, A. Payberah, J.R. Ayala Solares, F. Raimondi, M. Nazarzadeh, D. Canoy, K. Rahimi, Predicting the risk of emergency hospital admissions in the general population: development and validation of machine learning models in a cohort study using large-scale linked electronic health records, *PLOS Med.*, 15 (11).
- [20] P. Nguyen, T. Tran, N. Wickramasinghe, S. Venkatesh, Deepr: a convolutional net for medical records, *IEEE J. Biomed. Health Inform.* 21 (1) (2017) 22–30, <https://doi.org/10.1109/JBHI.2016.2633963>.
- [21] T. Pham, T. Tran, D. Phung, S. Venkatesh, Predicting healthcare trajectories from medical records: a deep learning approach, *J. Biomed. Inform.* 69 (2017) 218–229, <https://doi.org/10.1016/j.jbi.2017.04.001>.
- [22] A. Rajkomar, E. Oren, K. Chen, A.M. Dai, N. Hajaj, M. Hardt, P.J. Liu, X. Liu, J. Marcus, M. Sun, P. Sundberg, H. Yee, K. Zhang, Y. Zhang, G. Flores, G.E. Duggan, J. Irvine, Q. Le, K. Litsch, A. Mossin, J. Tansuwan, D. Wang, J. Wexler, J. Wilson, D. Ludwig, S.L. Volchenboum, K. Chou, M. Pearson, S. Madabushi, N.H. Shah, A.J. Butte, M.D. Howell, C. Cui, G.S. Corrado, J. Dean, Scalable and accurate deep learning with electronic health records, *npj Digital Med.* 1 (1) (2018) 18, <https://doi.org/10.1038/s41746-018-0029-1>.
- [23] Y. LeCun, Y. Bengio, G. Hinton, Deep learning, *Nature* 521 (2015) 436 EP, <https://doi.org/10.1038/nature14539>.
- [24] I. Goodfellow, Y. Bengio, A. Courville, *Deep Learning*, MIT Press, 2016 <http://www.deeplearningbook.org>.
- [25] J. Hirschberg, C.D. Manning, Advances in natural language processing, *Science* 349 (2015) 261–266.
- [26] A. Graves, Speech recognition with deep recurrent neural networks, 2013 IEEE International Conference on Acoustics, Speech and Signal Processing, 2013, pp. 6645–6649.
- [27] C. Szegedy, S. Ioffe, V. Vanhoucke, Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning, in: *AAAI*, 2017.
- [28] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, M.A. Riedmiller, Playing Atari with Deep Reinforcement Learning, *Computing Research Repository abs/1312.5602*.
- [29] D. Ravi, C. Wong, F. Deligianni, M. Berthelot, J. Andreu-Perez, B. Lo, G. Yang, Deep learning for health informatics, *IEEE J. Biomed. Health Inform.* 21 (1) (2017) 4–21, <https://doi.org/10.1109/JBHI.2016.2636665>.
- [30] D. Shen, G. Wu, H.-I. Suk, Deep learning in medical image analysis, *Annu. Rev. Biomed. Eng.* 19 (2017) 221–248.
- [31] V. Gulshan, L. Peng, M. Coram, M.C. Stumpe, D. Wu, A. Narayanaswamy, S. Venugopalan, K. Widner, T. Madams, J. Cuadros, et al., Development and validation of a deep learning algorithm for detection of diabetic retinopathy in retinal fundus photographs, *JAMA* 316 (22) (2016) 2402–2410.
- [32] Y. Liu, K.K. Gadepalli, M. Norouzi, G. Dahl, T. Kohlberger, S. Venugopalan, A.S. Boyko, A. Timofeev, P.Q. Nelson, G. Corrado, J. Hipp, L. Peng, M. Stumpe, Detecting Cancer Metastases on Gigapixel Pathology Images, *Tech. rep.*, arXiv (2017). URL <https://arxiv.org/abs/1703.02442>.
- [33] S. Webb, Deep learning for biology, *Nature*.
- [34] R. Poplin, P.-C. Chang, D. Alexander, S. Schwartz, T. Colthurst, A. Ku, D. Newburger, J. Djamco, N. Nguyen, P.T. Afshar, et al., A universal snp and small-indel variant caller using deep neural networks, *Nat. Biotechnol.*
- [35] C. Angermueller, T. Pärnamaa, L. Parts, O. Stegle, Deep learning for computational biology, *Molecular Systems Biology* 12 (7). <https://doi.org/10.15252/msb.20156651>.
- [36] T. Yue, H. Wang, Deep Learning for Genomics: A Concise Overview, arXiv.
- [37] E. Gawehn, J.A. Hiss, G. Schneider, Deep learning in drug discovery, *Mol. Inform.* 35 (1) (2016) 3–14.
- [38] R. Miotto, L. Li, B.A. Kidd, J.T. Dudley, Deep patient: an unsupervised representation to predict the future of patients from the electronic health records, *Scient. Rep.*, vol. 6.
- [39] J. Jameson, D.L. Longo, Precision medicine - personalized, problematic, and promising, *New Engl. J. Med.* 372 (2015) 2229–2234.
- [40] T. Tran, T.D. Nguyen, D. Phung, S. Venkatesh, Learning vector representation of medical objects via EMR-driven nonnegative restricted Boltzmann machines (eNRBM), *J. Biomed. Inform.* 54 (2015) 96–105, <https://doi.org/10.1016/j.jbi.2015.01.012>.
- [41] T. Mikolov, I. Sutskever, K. Chen, G.S. Corrado, J. Dean, Distributed representations of words and phrases and their compositionality, in: *Advances in Neural Information Processing Systems*, 2013, pp. 3111–3119.
- [42] E. Choi, M.T. Bahadori, E. Searles, C. Coffey, M. Thompson, J. Bost, J. Tejedro-Sojo, J. Sun, Multi-layer representation learning for medical concepts, *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ACM, 2016, pp. 1495–1504.
- [43] T. Mikolov, K. Chen, G.S. Corrado, J. Dean, Efficient Estimation of Word Representations in Vector Space, *Computing Research Repository*.
- [44] J. Pennington, R. Socher, C. Manning, Glove: Global Vectors for Word Representation, *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2014, pp. 1532–1543.
- [45] E. Choi, M.T. Bahadori, J. Sun, J. Kulas, A. Schuetz, W. Stewart, RETAIN: An Interpretable Predictive Model for Healthcare using Reverse Time Attention Mechanism, in: *Advances in Neural Information Processing Systems*, 2016, pp. 3504–3512.
- [46] E. Choi, M.T. Bahadori, A. Schuetz, W.F. Stewart, J. Sun, Doctor AI: Predicting Clinical Events via Recurrent Neural Networks, in: *Proceedings of the 1st Machine Learning for Healthcare Conference*, Vol. 56 of *Proceedings of Machine Learning Research*, PMLR, 2016, pp. 301–318. <<http://proceedings.mlr.press/v56/Choi16.html>>.
- [47] E. Choi, A. Schuetz, W.F. Stewart, J. Sun, Using recurrent neural network models for early detection of heart failure onset, *J. Am. Med. Inform. Assoc.* 24 (2) (2017) 361–370, <https://doi.org/10.1093/jamia/ocw112>.
- [48] B.C. Kwon, M.-J. Choi, J.T. Kim, E. Choi, Y.B. Kim, S. Kwon, J. Sun, J. Choo, RetainVis: visual analytics with interpretable and interactive recurrent neural networks on electronic medical records, *IEEE Trans. Visual. Comput. Graph.*
- [49] F. Ma, R. Chitta, J. Zhou, Q. You, T. Sun, J. Gao, Dipole: Diagnosis Prediction in Healthcare via Attention-based Bidirectional Recurrent Neural Networks, in: *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ACM, 2017, pp. 1903–1911.
- [50] E. Choi, M.T. Bahadori, L. Song, W.F. Stewart, J. Sun, GRAM: Graph-based Attention Model for Healthcare Representation Learning, in: *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '17, 2017, pp. 787–795.
- [51] J.T. Connor, R.D. Martin, L.E. Atlas, Recurrent neural networks and robust time series prediction, *IEEE Trans. Neural Networks* 5 (2) (1994) 240–254.
- [52] M. Schuster, K.K. Paliwal, Bidirectional recurrent neural networks, *IEEE Trans. Signal Process.* 45 (1997) 2673–2681.
- [53] S. Hochreiter, J. Schmidhuber, Long short-term memory, *Neural Comput.* 9 (1997) 1735–1780.
- [54] Z.C. Lipton, D.C. Kale, C. Elkan, R.C. Wetzel, Learning to Diagnose with LSTM Recurrent Neural Networks, *Computing Research Repository*.
- [55] L. Rokach, Ensemble-based classifiers, *Artif. Intell. Rev.* 33 (2009) 1–39.
- [56] T. Ma, C. Xiao, F. Wang, Health-ATM: A deep architecture for multifaceted patient health record representation and risk prediction, *Proceedings of the 2018 SIAM International Conference on Data Mining*, SIAM, 2018, pp. 261–269.
- [57] E. Herrett, A.M. Gallagher, K. Bhaskaran, H. Forbes, R. Mathur, T. van Staa, L. Smeeth, Data resource profile: clinical practice research Datalink (CPRD), *Int. J. Epidemiol.* 44 (3) (2015) 827–836.
- [58] Clinical Practice Research Datalink, <www.cprd.com> (Accessed: 11/September/2018).
- [59] Read Codes, <<https://digital.nhs.uk/services/terminology-and-classifications/read-codes>> (Accessed: 19/October/2018).
- [60] ICD-10 online versions, <<https://www.who.int/classifications/icd/icdonlineversions/en/>> (Accessed: 19/October/2018).
- [61] Clinical Classifications Software (CCS) for ICD-9-CM, <<https://www.hcup-us.ahrq.gov/toolssoftware/ccs/ccs.jsp>> (Accessed: 12/October/2018).
- [62] BNF Publications, <<https://www.bnf.org/>> (Accessed: 12/October/2018).
- [63] A. Stevenson, S. Deeny, R. Friebel, T. Gardner, R. Thorlby, Briefing: Emergency hospital admissions in England Which may be avoidable and how?, *Tech. rep.*, The Health Foundation (May 2018).
- [64] N. Conrad, A. Judge, J. Tran, H. Mohseni, D. Hedgecote, A.P. Crespillo, M. Allison, H. Hemingway, J.G. Cleland, J.J. McMurray, et al., Temporal trends and patterns in heart failure incidence: a population-based study of 4 million individuals, *The Lancet* 391 (10120) (2018) 572–580.
- [65] Reimplementation of RETAIN Recurrent Neural Network in Keras, <<https://github.com/Optom/retain-keras>> (Accessed: 11/October/2018).
- [66] A. Paszke, et al., PyTorch, <<https://pytorch.org/>> (2017).
- [67] F. Chollet et al., Keras, <<https://keras.io/>>, 2015.
- [68] L. Buitinck, G. Louppe, M. Blondel, F. Pedregosa, A. Mueller, O. Grisel, V. Niculae, P. Prettenhofer, A. Gramfort, J. Grobler, R. Layton, J. VanderPlas, A. Joly, B. Holt,

- G. Varoquaux, API design for machine learning software: experiences from the scikit-learn project, in: *EML PKDD Workshop: Languages for Data Mining and Machine Learning*, 2013, pp. 108–122.
- [69] M. Pelikan, D.E. Goldberg, E. Cantú-Paz, BOA: The Bayesian Optimization Algorithm, in: *Proceedings of the 1st Annual Conference on Genetic and Evolutionary Computation - Volume 1, GECCO'99*, Morgan Kaufmann Publishers Inc., 1999, pp. 525–532.
- [70] J. Snoek, H. Larochelle, R.P. Adams, Practical Bayesian Optimization of Machine Learning Algorithms, in: F. Pereira, C.J.C. Burges, L. Bottou, K.Q. Weinberger (Eds.), *Advances in Neural Information Processing Systems*, 2012, pp. 2951–2959.
- [71] skopt module, <<https://scikit-optimize.github.io/>> (Accessed: 12/September/2018).
- [72] T. Fawcett, ROC graphs: notes and practical considerations for researchers, *Mach. Learn.* 31 (1) (2004) 1–38.
- [73] T. Fawcett, An introduction to ROC analysis, *Pattern Recogn. Lett.* 27 (8) (2006) 861–874.
- [74] H. He, E.A. Garcia, Learning from imbalanced data, *IEEE Trans. Knowl. Data Eng.* 21 (2009) 1263–1284.
- [75] T. Saito, M. Rehmsmeier, The Precision-Recall Plot Is More Informative than the ROC Plot When Evaluating Binary Classifiers on Imbalanced Datasets, *PLoS One* 10 (3) (2015) e0118432.
- [76] Precision-Recall (scikit-learn), <http://scikit-learn.org/stable/auto_examples/model_selection/plot_precision_recall.html> (Accessed: 05/November/2018).
- [77] Average Precision score (scikit-learn), <http://scikit-learn.org/stable/modules/generated/sklearn.metrics.average_precision_score.html> (Accessed: 05/November/2018).
- [78] Y. Yuan, W. Su, M. Zhu, Threshold-free measures for assessing the performance of medical screening tests, *Front. Public Health* 3 (2015) 57.
- [79] K. Boyd, K.H. Eng, C.D. Page, Area Under the Precision-Recall Curve: Point Estimates and Confidence Intervals, *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, Springer, 2013, pp. 451–466.
- [80] F1-Score (scikit-learn), <https://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html> (Accessed: 01/June/2019).
- [81] C.A. Emdin, S.G. Anderson, T. Callender, N. Conrad, G. Salimi-Khorshidi, H. Mohseni, M. Woodward, K. Rahimi, Usual blood pressure, peripheral arterial disease, and vascular risk: cohort study of 4.2 million adults, *Br. Med. J.* 351 (2015) h4865.
- [82] C.A. Emdin, P.M. Rothwell, G. Salimi-Khorshidi, A. Kiran, N. Conrad, T. Callender, Z. Mehta, S.T. Pendlebury, S.G. Anderson, H. Mohseni, et al., Blood pressure and risk of vascular dementia: evidence from a primary care registry and a cohort study of transient ischemic attack and stroke, *Stroke* 47 (6) (2016) 1429–1435.
- [83] K. Rahimi, H. Mohseni, C.M. Otto, N. Conrad, J. Tran, M. Nazarzadeh, M. Woodward, T. Dwyer, S. MacMahon, Elevated blood pressure and risk of mitral regurgitation: A longitudinal cohort study of 5.5 million United Kingdom adults, *PLoS Med.*, 14 (10).
- [84] CPRD Bibliography, <<https://www.cprd.com/bibliography>> (Accessed: 17/December/2018).
- [85] J. Davis, M. Goadrich, The relationship between precision-recall and roc curves, *Proceedings of the 23rd International Conference on Machine learning*, ACM, 2006, pp. 233–240.
- [86] N.V. Chawla, K.W. Bowyer, L.O. Hall, W.P. Kegelmeyer, SMOTE: synthetic minority over-sampling technique, *J. Artif. Intell. Res.* 16 (2002) 321–357.
- [87] G. Lemaitre, F. Nogueira, C.K. Aridas, Imbalanced-learn: A Python toolbox to tackle the curse of imbalanced datasets in machine learning, *J. Mach. Learn. Res.* 18 (17) (2017) 1–5 <http://jmlr.org/papers/v18/16-365>.
- [88] L. McInnes, J. Healy, UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction, *arXiv preprint arXiv:1802.03426*.
- [89] Understanding the High Prevalence of Low Prevalence Chronic Disease Combinations: Databases and Methods for Research, <https://goo.gl/srsZs2>, 4/December/2018.
- [90] WHO, Introduction to Drug Utilization Research, World Health Organization, 2003.
- [91] J. Snoek, O. Rippel, K. Swersky, R. Kiros, N. Satish, N. Sundaram, M.M.A. Patwary, Prabhakar, R.P. Adams, Scalable Bayesian Optimization Using Deep Neural Networks, in: *ICML*, 2015.
- [92] J.T. Springenberg, A. Klein, S. Falkner, F. Hutter, Bayesian Optimization with Robust Bayesian Neural Networks, in: *NIPS*, 2016.
- [93] R. Miikkilainen, J.Z. Liang, E. Meyerson, A. Rawal, D. Fink, O. Francon, B. Raju, H. Shahrzad, A. Navruzian, N. Duffy, B. Hodjat, Evolving Deep Neural Networks, *Computing Research Repository abs/1703.00548*.
- [94] J.K. Dutta, J. Liu, U. Kurup, M. Shah, Effective Building Block Design for Deep Convolutional Neural Networks using Search, *Computing Research Repository abs/1801.08577*.
- [95] J. Howard et al., fastai, <<https://github.com/fastai/fastai>>, 2018.
- [96] L.N. Smith, Cyclical Learning Rates for Training Neural Networks, 2017 *IEEE Winter Conference on Applications of Computer Vision (WACV)*, 2017, pp. 464–472.
- [97] I. Loshchilov, F. Hutter, SGDR: Stochastic Gradient Descent with Warm Restarts, 2016.
- [98] L. Jacobs, L. Efremov, J.P. Ferreira, L. Thijs, W.-Y. Yang, Z.-Y. Zhang, R. Latini, S. Masson, N. Agabiti, P. Sever, et al., Risk for incident heart failure: a subject-level meta-analysis from the heart OMics in AGEing (HOMAGE) Study, *J. Am. Heart Assoc.* 6 (5) (2017) e005231.
- [99] Y. Gal, Z. Ghahramani, Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning, in: *Proceedings of The 33rd International Conference on Machine Learning*, Vol. 48 of *Proceedings of Machine Learning Research*, New York, New York, USA, 2016, pp. 1050–1059.
- [100] Y. Gal, Uncertainty in Deep Learning, University of Cambridge.
- [101] D. Tran, D. Mike, M. van der Wilk, D. Hafner, Bayesian Layers: A Module for Neural Network Uncertainty, *arXiv preprint arXiv:1812.03973*.
- [102] S. Chakraborty, R. Tomsett, R. Raghavendra, D. Harborne, M. Alzantot, F. Cerutti, M. Srivastava, A. Preece, S. Julier, R.M. Rao, T.D. Kelley, D. Braines, M. Sensoy, C. J. Willis, P. Gurram, Interpretability of deep learning models: A survey of results, in: *2017 IEEE SmartWorld, Ubiquitous Intelligence Computing, Advanced Trusted Computing, Scalable Computing Communications, Cloud Big Data Computing, Internet of People and Smart City Innovation (SmartWorld/SCALCOM/UIC/ATC/CBDCom/IOP/SCI)*, 2017, pp. 1–6.
- [103] H.J. Kam, H.Y. Kim, Learning representations for the early detection of sepsis with deep neural networks, *Comput. Biol. Med.* 89 (2017) 248–255.
- [104] B.K. Reddy, D. Delen, Predicting hospital readmission for lupus patients: an RNN-LSTM-based deep-learning methodology, *Comput. Biol. Med.* 101 (2018) 199–209.
- [105] J. Chu, W. Dong, K. He, H. Duan, Z. Huang, Using neural attention networks to detect adverse medical events from electronic health records, *J. Biomed. Inform.*
- [106] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A.N. Gomez, L. Kaiser, I. Polosukhin, Attention Is All You Need, in: *NIPS*, 2017.
- [107] J. Devlin, M.-W. Chang, K. Lee, K. Toutanova, BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding, *Computing Research Repository abs/1810.04805*.
- [108] M.E. Peters, M. Neumann, M. Iyyer, M. Gardner, C. Clark, K. Lee, L. Zettlemoyer, Deep contextualized word representations, in: *Proc. of NAACL*, 2018.
- [109] S. Ruder, J. Howard, Universal Language Model Fine-tuning for Text Classification, in: *ACL*, 2018.
- [110] M. Zeng, M. Li, Z. Fei, Y. Yu, Y. Pan, J. Wang, Automatic ICD-9 coding via deep transfer learning, *Neurocomputing*.
- [111] Y. Wang, S. Liu, N. Afzal, M. Rastegar-Mojarad, L. Wang, F. Shen, P. Kingsbury, H. Liu, A comparison of word embeddings for the biomedical natural language processing, *J. Biomed. Inform.* 87 (2018) 12–20.
- [112] B. Lim, M. van der Schaar, Disease-Atlas: Navigating Disease Trajectories with Deep Learning, *Comput. Res. Reposit., abs/1803.10254*.
- [113] S. Lee, Natural language generation for electronic health records, *Comput. Res. Reposit., abs/1806.01353*.